

Phase 1: Foundation Setup (Weeks 1-6)

Overview

Phase 1 establishes the core infrastructure and foundational components required for the BioSpark Health AI system. This phase focuses on environment setup, basic security, and initial API framework.

Week-by-Week Implementation Plan

Week 1: Environment Setup

Objectives: Establish development and staging environments

Tasks:

1. Infrastructure Provisioning

```
```bash
Kubernetes cluster setup
kubectl create namespace biospark-dev
kubectl create namespace biospark-staging

Install essential operators
helm install ingress-nginx ingress-nginx/ingress-nginx
helm install cert-manager jetstack/cert-manager
```
```

1. Database Setup

```
sql
-- PostgreSQL setup
CREATE DATABASE biospark_dev;
CREATE DATABASE biospark_staging;
CREATE USER biospark_user WITH PASSWORD 'secure_password';
GRANT ALL PRIVILEGES ON DATABASE biospark_dev TO biospark_user;
```

2. Container Registry Configuration

```
yaml
# docker-registry-secret.yaml
apiVersion: v1
kind: Secret
metadata:
  name: docker-registry-secret
type: kubernetes.io/dockerconfigjson
data:
  .dockerconfigjson: <base64-encoded-docker-config>
```

Deliverables:

- ☒ Kubernetes clusters operational
- ☒ Database instances configured
- ☒ Container registry accessible
- ☒ Basic networking configured

Week 2: Core Infrastructure

Objectives: Deploy foundational services and monitoring

Tasks:

1. Service Mesh Deployment

```
bash
# Istio installation
istioctl install --set values.defaultRevision=default
kubectl label namespace biospark-dev istio-injection=enabled
```

1. Monitoring Stack

```
yaml
# prometheus-values.yaml
prometheus:
  prometheusSpec:
    retention: 30d
    storageSpec:
      volumeClaimTemplate:
        spec:
          storageClassName: fast-ssd
          resources:
            requests:
              storage: 100Gi
```

2. Logging Infrastructure

```
bash
# ELK Stack deployment
helm install elasticsearch elastic/elasticsearch
helm install kibana elastic/kibana
helm install filebeat elastic/filebeat
```

Deliverables:

- ☒ Service mesh operational
- ☒ Monitoring stack deployed
- ☒ Logging infrastructure ready
- ☒ Basic alerting configured

Week 3: Security Foundation

Objectives: Implement core security components

Tasks:

1. Secrets Management

```
bash
# Vault deployment
helm install vault hashicorp/vault
vault auth enable kubernetes
vault policy write biospark-policy - <<EOF
path "secret/data/biospark/*" {
  capabilities = ["read"]
```

```
}
```

```
EOF
```

1. Certificate Management

```
yaml
# cluster-issuer.yaml
apiVersion: cert-manager.io/v1
kind: ClusterIssuer
metadata:
  name: letsencrypt-prod
spec:
  acme:
    server: https://acme-v02.api.letsencrypt.org/directory
    email: admin@biospark.ai
    privateKeySecretRef:
      name: letsencrypt-prod
```

2. Network Policies

```
```yaml
network-policy.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
 name: biospark-network-policy
spec:
 podSelector:
 matchLabels:
 app: biospark
 policyTypes:
 - Ingress
 - Egress
```
```

Deliverables:

- ☒ Secrets management operational
- ☒ TLS certificates automated
- ☒ Network security policies active
- ☒ Basic RBAC configured

Week 4: API Gateway Setup

Objectives: Deploy and configure API gateway

Tasks:

1. Kong Gateway Deployment

```
yaml
# kong-values.yaml
image:
  repository: kong
  tag: "3.4"
env:
  database: postgres
```

```
pg_host: postgresql
pg_database: kong
proxy:
  type: LoadBalancer
```

1. Rate Limiting Configuration

```
bash
# Configure rate limiting plugin
curl -X POST http://kong-admin:8001/plugins \
  --data "name=rate-limiting" \
  --data "config.minute=1000" \
  --data "config.policy=local"
```

2. Authentication Plugin

```
bash
# JWT authentication setup
curl -X POST http://kong-admin:8001/plugins \
  --data "name=jwt" \
  --data "config.secret_is_base64=false"
```

Deliverables:

- ☒ API Gateway deployed
- ☒ Rate limiting configured
- ☒ Authentication plugins active
- ☒ Basic routing rules defined

Week 5: Database Schema Implementation

Objectives: Create core database schema and initial data structures

Tasks:

1. Core Tables Creation

```
```sql
- patients table
CREATE TABLE patients (
 id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
 first_name VARCHAR(100) NOT NULL,
 last_name VARCHAR(100) NOT NULL,
 date_of_birth DATE NOT NULL,
 email VARCHAR(255) UNIQUE,
 phone VARCHAR(20),
 created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
 updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

- health_records table
CREATE TABLE health_records (
 id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
 patient_id UUID REFERENCES patients(id),
 record_type VARCHAR(50) NOT NULL,
 data JSONB NOT NULL,
 recorded_at TIMESTAMP NOT NULL,
```

```

created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

-- agents table
CREATE TABLE agents (
id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
name VARCHAR(100) NOT NULL,
type VARCHAR(50) NOT NULL,
config JSONB,
status VARCHAR(20) DEFAULT 'inactive',
created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

```

### 1. Indexes and Constraints

```

-- sql
-- Performance indexes
CREATE INDEX idx_patients_email ON patients(email);
CREATE INDEX idx_health_records_patient_id ON health_records(patient_id);
CREATE INDEX idx_health_records_recorded_at ON health_records(recorded_at);
CREATE INDEX idx_agents_type ON agents(type);
CREATE INDEX idx_agents_status ON agents(status);

-- JSONB indexes for health data
CREATE INDEX idx_health_records_data_gin ON health_records USING GIN(data);

```





### 1. Initial Data Migration

```

sql
-- Insert default agent configurations
INSERT INTO agents (name, type, config, status) VALUES
('Analyst Agent', 'analyst', '{"version": "1.0", "capabilities": ["data_analysis", "reporting"]}', 'active'),
('Architect Agent', 'architect', '{"version": "1.0", "capabilities": ["system_design", "optimization"]}', 'active'),
('Developer Agent', 'developer', '{"version": "1.0", "capabilities": ["code_generation", "testing"]}', 'active'),
('Orchestrator Agent', 'orchestrator', '{"version": "1.0", "capabilities": ["workflow_management", "coordination"]}', 'active');

```

#### Deliverables:

-  Core database schema implemented
-  Performance indexes created
-  Initial data populated
-  Database migrations framework ready

## Week 6: Basic API Framework

**Objectives:** Implement foundational API endpoints and authentication

#### Tasks:

##### 1. Authentication Service

```

```typescript
// auth.service.ts
import jwt from 'jsonwebtoken';
import bcrypt from 'bcrypt';

export class AuthService {
  async login(email: string, password: string): Promise {
    const user = await this.validateUser(email, password);
    if (!user) throw new Error('Invalid credentials');

```

```

      return jwt.sign(
        { userId: user.id, email: user.email },
        process.env.JWT_SECRET!,
        { expiresIn: '24h' }
      );
    }

    async validateToken(token: string): Promise<any> {
      return jwt.verify(token, process.env.JWT_SECRET!);
    }

```

```

}
```

```

## 1. Core API Endpoints

```

```typescript
// patients.controller.ts
import { Request, Response } from 'express';
import { PatientService } from '../services/patient.service';

export class PatientsController {
  private patientService = new PatientService();

```

```

  async getPatients(req: Request, res: Response) {
    const { limit = 50, offset = 0 } = req.query;
    const patients = await this.patientService.findAll({
      limit: Number(limit),
      offset: Number(offset)
    });
    res.json(patients);
  }

  async createPatient(req: Request, res: Response) {
    const patient = await this.patientService.create(req.body);
    res.status(201).json(patient);
  }

```

```

}
```

```

## 1. API Documentation

```

```yaml
# openapi.yaml
openapi: 3.0.0
info:

```

title: BioSpark Health AI API
version: 1.0.0
description: Core API for BioSpark Health AI system

paths:

/api/v1/auth/login:

post:

summary: User authentication

requestBody:

required: true

content:

application/json:

schema:

type: object

properties:

email:

type: string

password:

type: string

responses:

200:

description: Authentication successful

content:

application/json:

schema:

type: object

properties:

token:

type: string

...

Deliverables:

- ☒ Authentication service operational
- ☒ Core API endpoints implemented
- ☒ API documentation generated
- ☒ Basic error handling configured

Phase 1 Quality Gates

Technical Validation

- ☐ All infrastructure components deployed successfully
- ☐ Database schema matches specifications
- ☐ API endpoints respond within 200ms
- ☐ Authentication system functional
- ☐ Basic security measures active

Security Review

- ☐ Secrets properly managed in Vault
- ☐ TLS certificates valid and auto-renewing
- ☐ Network policies restricting traffic

- [] RBAC permissions configured correctly
- [] No hardcoded credentials in code

Performance Benchmarks

- [] API Gateway handling 1000 req/min
- [] Database queries under 50ms
- [] System startup time under 5 minutes
- [] Memory usage within allocated limits
- [] CPU utilization under 70%

Documentation Requirements

- [] Infrastructure setup documented
- [] API documentation complete
- [] Security procedures documented
- [] Troubleshooting guides available
- [] Runbooks for common operations

Troubleshooting Guide

Common Issues

Database Connection Issues

```
# Check database connectivity
kubectl exec -it postgres-pod -- psql -U biospark_user -d biospark_dev -c "SELECT 1;"

# Verify secrets
kubectl get secret postgres-secret -o yaml
```

API Gateway Not Responding

```
# Check Kong status
kubectl get pods -l app=kong
kubectl logs -l app=kong

# Verify service configuration
kubectl get services kong-proxy
```

Certificate Issues

```
# Check certificate status
kubectl get certificates
kubectl describe certificate biospark-tls

# Force certificate renewal
kubectl delete certificate biospark-tls
```

Phase 1 establishes the solid foundation required for successful implementation of subsequent phases.